

Basic Introduction to gcloud

Nick Thompson

January 18, 2016

Goals of this discussion

Goals of this discussion

- ▶ Learn how to rent CPUs and RAM from google

Goals of this discussion

- ▶ Learn how to rent CPUs and RAM from google
- ▶ Learn how to rent persistent disks

Goals of this discussion

- ▶ Learn how to rent CPUs and RAM from google
- ▶ Learn how to rent persistent disks
- ▶ Learn how to control expenses

Goals of this discussion

- ▶ Learn how to rent CPUs and RAM from google
- ▶ Learn how to rent persistent disks
- ▶ Learn how to control expenses
- ▶ Learn how to manage permissions

Goals of this discussion

- ▶ Learn how to rent CPUs and RAM from google
- ▶ Learn how to rent persistent disks
- ▶ Learn how to control expenses
- ▶ Learn how to manage permissions
- ▶ Learn how to create teams of identical VMs

Goals of this discussion

- ▶ Learn how to rent CPUs and RAM from google
- ▶ Learn how to rent persistent disks
- ▶ Learn how to control expenses
- ▶ Learn how to manage permissions
- ▶ Learn how to create teams of identical VMs
- ▶ Learn how to autoscale nodes based on load

What we won't discuss, but should

What we won't discuss, but should

- ▶ The Python/Go/Node bindings

What we won't discuss, but should

- ▶ The Python/Go/Node bindings
- ▶ The REST (representational state transfer) interface

What we won't discuss, but should

- ▶ The Python/Go/Node bindings
- ▶ The REST (representational state transfer) interface
- ▶ gcloud sql

What we won't discuss, but should

- ▶ The Python/Go/Node bindings
- ▶ The REST (representational state transfer) interface
- ▶ gcloud sql
- ▶ Google Container Engine

What we won't discuss, but should

- ▶ The Python/Go/Node bindings
- ▶ The REST (representational state transfer) interface
- ▶ gcloud sql
- ▶ Google Container Engine
- ▶ Lots more . . .

Getting started:

```
$ curl https://sdk.cloud.google.com | bash
$ exec -l $SHELL
$ gcloud init
$ gcloud auth login
$ firefox https://console.developers.google.com
```

You don't need a gmail account to use gcloud, but you must enter an email into gcloud that then becomes your google account login.

Starting a Project

Filter by name, ID, or label

Project ID
graphical-calm-976

and pending deletion

Columns ▾

Labels

New Project

Project name ?

Your project ID will be fiery-buttress-118421 [Edit](#)

[Show advanced options...](#)

Create

Cancel

Starting a Project

Each project has

Starting a Project

Each project has

- ▶ its bill separated in the invoice (1 billing account \mapsto many project accounts)

Starting a Project

Each project has

- ▶ its bill separated in the invoice (1 billing account \mapsto many project accounts)
- ▶ its own VMs

Starting a Project

Each project has

- ▶ its bill separated in the invoice (1 billing account \mapsto many project accounts)
- ▶ its own VMs
- ▶ its own persistent disks

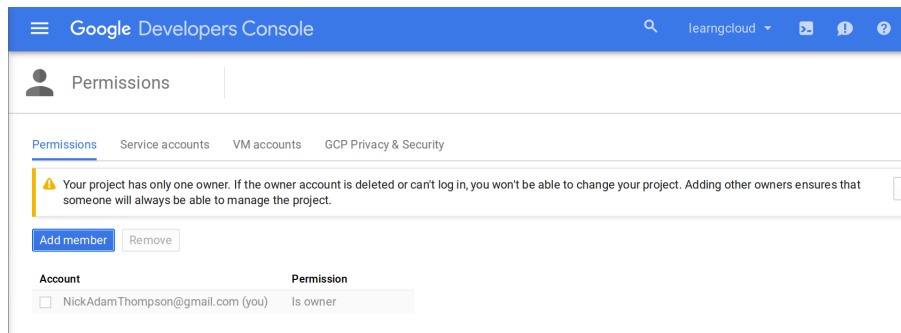
Starting a Project

Each project has

- ▶ its bill separated in the invoice (1 billing account \mapsto many project accounts)
- ▶ its own VMs
- ▶ its own persistent disks
- ▶ its own users with their associated permissions.

Starting a Project

Multiple people can be listed as project owners, or given read access, or read/write access to the project:



The screenshot shows the Google Developers Console interface. The top navigation bar is blue with the text "Google Developers Console" and a search icon. Below the navigation bar, the "Permissions" section is active, indicated by a person icon and the word "Permissions". Underneath, there are tabs for "Permissions", "Service accounts", "VM accounts", and "GCP Privacy & Security". A warning message with a yellow triangle icon states: "Your project has only one owner. If the owner account is deleted or can't log in, you won't be able to change your project. Adding other owners ensures that someone will always be able to manage the project." Below the warning, there are two buttons: "Add member" (in blue) and "Remove" (in grey). At the bottom, there is a table with two columns: "Account" and "Permission". The table contains one entry: a checkbox, the email "NickAdamThompson@gmail.com (you)", and the role "Is owner".

Account	Permission
<input type="checkbox"/> NickAdamThompson@gmail.com (you)	Is owner

Note that people who only have read access to the project nonetheless have root access to all the VMs!

Permissions

The permission levels for a project are

Permissions

The permission levels for a project are

- ▶ Owners—who can add/remove team members, and rent resources, and are root on resources

Permissions

The permission levels for a project are

- ▶ Owners—who can add/remove team members, and rent resources, and are root on resources
- ▶ Editors—who can rent resources, and are root on resources

Permissions

The permission levels for a project are

- ▶ Owners—who can add/remove team members, and rent resources, and are root on resources
- ▶ Editors—who can rent resources, and are root on resources
- ▶ Viewers—who can't rent resources, but are root on resources

Starting a Project

When you grant someone permission to view/edit/co-own your project, they receive an email asking them if they want to join:

Hello, 张洁,

I invite you to join the Google Developers Console project "myfirstproject". Please click this link to accept my invitation:

<https://console.developers.google.com/project/learncloud-1184/rsvp?account=zjzjhn@gmail.com>

Thanks,

Nick Thompson

NickAdamThompson@gmail.com



Google Cloud Platform

Build your apps on Google's infrastructure.

Take advantage of the power and performance of Google APIs like Maps, Social, Translate, Android, and much more. [Learn more](#)

In addition, Google Cloud Platform lets you build, deploy, and scale applications, websites, and services on the same infrastructure that runs Google. [Learn more](#)

© 2014 Google Inc. 1600 Amphitheatre Parkway, Mountain View, CA 94043

gcloud user accounts

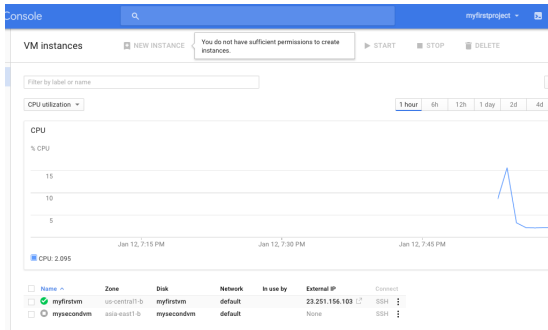
gcloud user accounts are in beta, and seem to be evolving rapidly. To see more, use

```
$ gcloud beta compute users -h
```

or visit the [docs](#).

gcloud user accounts

Users with read-only access to your project have root on your VMs, but they can't launch new VMs on your dime:



Starting a project

You need tell the gcloud command-line tool that you've created a new project:

```
$ gcloud config list
[core]
account = nickadamthompson@gmail.com
disable_usage_reporting = True
project = graphical-cairn-97618
[meta]
active_config = default
```

If the project field has the wrong value, you need to set it:

```
$ gcloud config set project learncloud-1184
```

Note that you don't set the project *name*, you set the project *ID*.

Setting up an instance

Google Developers Console

Compute Engine

VM instances

Instance groups

Instance templates

Disks

Snapshots

Images

Metadata

Health checks

Zones

Operations

Quotas

Settings

Create a new instance

Name

instance-1

Zone

us-central1-c

Machine type

1 vCPU

3.75 GB memory

Customize

Boot disk

New 10 GB standard persistent disk

Image

Debian GNU/Linux 8.2 (jessie)

Change

Firewall

Add tags and firewall rules to allow specific network traffic from the Internet

☐ Allow HTTP traffic

☐ Allow HTTPS traffic

Project access

☐ Allow API access to all Google Cloud services in the same project. [Learn more](#)

Management

Disks

Networking

Access & security

Description (Optional)

Tags (Optional)

\$25.95 per month estimated

Effective hourly rate \$0.036 (730 hours per month)

[Details](#)

Setting up an instance

Things to choose at this point:

Setting up an instance

Things to choose at this point:

- ▶ Number of cores (min 1/2, max 32)

Setting up an instance

Things to choose at this point:

- ▶ Number of cores (min 1/2, max 32)
- ▶ amount of RAM (min 0.6GB, max unlimited, it seems, but 200GB is supported for sure)

Setting up an instance

Things to choose at this point:

- ▶ Number of cores (min 1/2, max 32)
- ▶ amount of RAM (min 0.6GB, max unlimited, it seems, but 200GB is supported for sure)
- ▶ size of disk, SSD (max 10TB) or spinning (max 10TB)

Setting up an instance

Things to choose at this point:

- ▶ Number of cores (min 1/2, max 32)
- ▶ amount of RAM (min 0.6GB, max unlimited, it seems, but 200GB is supported for sure)
- ▶ size of disk, SSD (max 10TB) or spinning (max 10TB)
- ▶ Operating system (Ubuntu, Centos, CoreOS), or choose a VM snapshot

Setting up an instance

Things to choose at this point:

- ▶ Number of cores (min 1/2, max 32)
- ▶ amount of RAM (min 0.6GB, max unlimited, it seems, but 200GB is supported for sure)
- ▶ size of disk, SSD (max 10TB) or spinning (max 10TB)
- ▶ Operating system (Ubuntu, Centos, CoreOS), or choose a VM snapshot
- ▶ Firewall rules

Setting up an instance

Things to choose at this point:

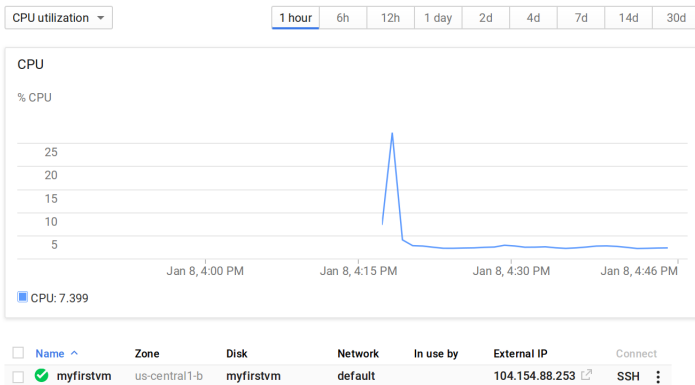
- ▶ Number of cores (min 1/2, max 32)
- ▶ amount of RAM (min 0.6GB, max unlimited, it seems, but 200GB is supported for sure)
- ▶ size of disk, SSD (max 10TB) or spinning (max 10TB)
- ▶ Operating system (Ubuntu, Centos, CoreOS), or choose a VM snapshot
- ▶ Firewall rules
- ▶ Whether to use static or ephemeral IP addresses (static IPs cost money!)

Setting up an instance

Aside: If you need more than 10TB of disk space, you need to fill out the [Google Compute Engine Quota Change Request Form](#).

Setting up an instance

Once you create a VM, you'll be assigned an external IP address and can see the load on your server:



Setting up an instance

To access your VM, use:

```
$ gcloud compute ssh myfirstvm --zone us-central1-b
```

(Don't choose the wrong zone or else your instance won't be found!)

This is really a wrapper script around the IP address of your instance:

```
$ ssh -i ~/.ssh/google_compute_engine 104.154.88.253
```

Setting up an instance

For every console action, there is a equivalent gcloud command. So, for example, to set up an instance, you could type

```
$ gcloud compute instances create mysecondvm \
    --image ubuntu-15-10 --zone us-central1-b
```

This is useful for scripting.

Exercise

Why doesn't the following code work?

```
local$ gcloud compute instances create vm1 --zone us-central1-b
local$ gcloud compute ssh vm1
vm1$ gcloud compute instances create vm2 --zone us-central1-b
ERROR: (gcloud.compute.instances.create) Failed to find image for alias
- Insufficient Permission
```

Solution

You are logged into gcloud, and can rent resources. However, your VM isn't you, and it doesn't have permission to rent resources on your behalf. You can solve this problem by authorizing your VM to be able to rent resources using a *service account*.

Service Accounts

To authorize a VM to rent resources on your behalf, use the *scopes* tag:

```
$ gcloud compute instances create vm-scoped \
  --scopes compute-rw --zone us-central1-b
```

This gives your vm permission to spawn new vms.

Startup Scripts

To get an image configured for autoscaling, we need to specify what sort of software we need on the image before it's creation. We can do this via a "startup script":

Automation

Startup script (Optional)

You can choose to specify a startup script that will run when your instance boots up or restarts. Start up scripts can be used to install software and updates, and to ensure that services are running within the virtual machine. [Learn more](#)

```
#!/bin/bash
sudo apt-get update
sudo apt-get install -y nginx
sudo chmod a+rw /var/www/html/index.nginx-debian.html
sudo echo "Our startup script works" > /var/www/html/index.nginx-debian.html
|
```

Startup Scripts

It works!

```
$ curl 130.211.135.33
```

Our startup script works

However, this is not a good design; as the startup script is not in source control.

Startup Scripts

Once your startup script is in source control, you can easily deploy a new instance via:

```
$ gcloud compute instances create wstartup \
  --metadata-from-file startup-script=script.sh \
  --zone=us-central1-b --tags "http-server"
```

Created [<https://www.googleapis.com/compute/v1/projects/learnngcloud-118>]

NAME	ZONE	MACHINE_TYPE	PREEMPTIBLE	INTERNAL_IP	EXTERNAL_I
wstartup	us-central1-a	n1-standard-1		10.240.0.5	104.197.59

```
$ curl 104.197.59.116
```

Our startup script works

Startup Scripts

Note that you can ssh into your machine before the startup script has finished!
To see if your startup script has finished, or to debug your startup script:

```
$ gcloud compute ssh wstartup --zone=us-central1-b
```

```
wstartup$ cat /var/log/startupscript.log
```

```
...
```

```
ubuntu startupscript: Finished running startup script /var/run/google.s
```

Regions and Zones

Google divides the world into regions:

```
$ gcloud compute regions list
```

NAME	CPUS	DISKS_GB	ADDRESSES	RESERVED_ADDRESSES	ST
asia-east1	0.00/24.00	10/10240	0/23	0/7	UP
europa-west1	0.00/24.00	0/10240	0/23	0/7	UP
us-central1	4.00/24.00	40/10240	4/23	0/7	UP
us-east1	0.00/24.00	0/10240	0/23	0/7	UP

Regions and Zones

And each region is divided into multiple zones

```
$ gcloud compute zones list
```

NAME	REGION	STATUS	NEXT_MAINTENANCE	TURNDOWN_DATE
asia-east1-a	asia-east1	UP		
asia-east1-b	asia-east1	UP		
asia-east1-c	asia-east1	UP		
europa-west1-b	europa-west1	UP		
europa-west1-d	europa-west1	UP		
europa-west1-c	europa-west1	UP		
us-central1-c	us-central1	UP		
us-central1-a	us-central1	UP		
us-central1-f	us-central1	UP		
us-central1-b	us-central1	UP		
us-east1-c	us-east1	UP		
us-east1-b	us-east1	UP		
us-east1-d	us-east1	UP		

Regions and Zones

A zone is essentially single datacenter, so two instances in a zone can communicate very quickly with one another:

```
vmcentral1b-1$ ping vmcentral1b-2 # ping VM in same datacenter/zone:
```

```
rtt min/avg/max/mdev = 0.308/0.394/0.841/0.104 ms
```

```
vmcentral1b-1$ ping vmcentral1a # Ping VM in same region, different zone
```

```
rtt min/avg/max/mdev = 0.571/0.673/1.210/0.099 ms
```

```
vmcentral1b-1$ ping vm-in-asia
```

```
rtt min/avg/max/mdev = 154.738/154.878/155.326/0.442 ms
```

So within-zone communication is fastest, within-region is fast, between-region is slow. (Note how your instance names are resolved via DNS!)

Regions and Zones

Things you rent from Google are classified by their scope in the global/regional/zonal hierarchy. For instance,

Regions and Zones

Things you rent from Google are classified by their scope in the global/regional/zonal hierarchy. For instance,

- ▶ Images, VM snapshots, firewall rules, and buckets are global resources

Regions and Zones

Things you rent from Google are classified by their scope in the global/regional/zonal hierarchy. For instance,

- ▶ Images, VM snapshots, firewall rules, and buckets are global resources
- ▶ Addresses are regional resources

Regions and Zones

Things you rent from Google are classified by their scope in the global/regional/zonal hierarchy. For instance,

- ▶ Images, VM snapshots, firewall rules, and buckets are global resources
- ▶ Addresses are regional resources
- ▶ Instances and their boot disks are zonal resources

Exercise

Why does google need to ask us for the zone for almost every command?

```
$ gcloud compute instances describe vm2
```

For the following instances:

- [vm2]

choose a zone:

```
[1] asia-east1-a
```

```
...
```

Solution

Everything action in gcloud is achieved via REST (representational state transfer).

This is a POST/GET/DELETE/PUT request to an html endpoint, and the endpoint url contains the zone!

Equivalent REST request

This is the REST request with the parameters you have selected.

```
POST https://www.googleapis.com/compute/v1/projects/learncloud-1184/zones/us-central1-b
{
  "name": "instance-1",
  "zone": "projects/learncloud-1184/zones/us-central1-b",
  "machineType": "projects/learncloud-1184/zones/us-central1-b/machineTypes/n1-standard-1",
  "metadata": {
    "items": []
  },
  "tags": {
    "items": []
  },
  "disks": [
    {
      "type": "PERSISTENT",
      "boot": true,
      "mode": "READ_WRITE",
      "autoDelete": true,
      "deviceName": "instance-1",
      "initializeParams": {
        "sourceImage": "https://www.googleapis.com/compute/v1/projects/debian-cloud/global/images/debian-9-stretch-v20170816",
        "diskType": "https://www.googleapis.com/compute/v1/projects/learncloud-1184/zones/us-central1-b/diskTypes/pd-standard",
        "diskSizeGb": "10"
      }
    }
  ]
}
```

☐ Line wrapping

Close

[REST API reference](#)

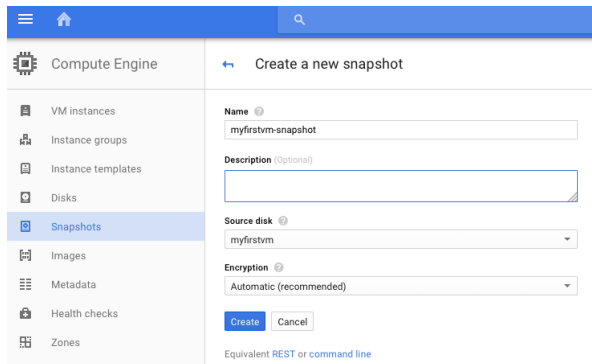
Snapshotting an instance

Once you have installed your favorite software on your VM, you might want to snapshot it:

```
myfirstvm$ sudo apt-get update
myfirstvm$ sudo apt-get install -y gcc g++ emacs git
myfirstvm$ sudo touch /opt/hello.txt && sudo chmod a+rw /opt/hello.txt
myfirstvm$ echo "Hello from GCE!" >> /opt/hello.txt
myfirstvm$ exit
localhost$ gcloud compute disks snapshot "myfirstvm" \
           --zone "us-central1-b" --snapshot-names "firstvmsnapshot"
```

Snapshotting an instance

Of course, using the console GUI is a bit easier:



The screenshot shows the Google Cloud Platform console interface for creating a new snapshot. The left sidebar contains a navigation menu with the following items: Compute Engine, VM instances, Instance groups, Instance templates, Disks, **Snapshots** (highlighted), Images, Metadata, Health checks, and Zones. The main content area is titled 'Create a new snapshot' and contains the following form fields:

- Name**: A text input field containing 'myfirstvm-snapshot'.
- Description**: A text input field with the label '(Optional)'.
- Source disk**: A dropdown menu showing 'myfirstvm'.
- Encryption**: A dropdown menu showing 'Automatic (recommended)'.

At the bottom of the form are two buttons: 'Create' and 'Cancel'. Below the buttons, there is a link: 'Equivalent [REST](#) or [command line](#)'.

Snapshotting an instance

Once you have a snapshot, you can create a boot disk from it:

Compute Engine

VM instances

Instance groups

Instance templates

Disks

Snapshots

Images

Metadata

Health checks

Zones

Operations

Quotas

Settings

Create a new disk

Name

firstvm-boot-disk

Description (Optional)

Zone

us-central1-b

Disk Type

Standard persistent disk

Source type

Image

Snapshot

None (blank disk)

Source snapshot

firstvmsnapshot

Size (GB) (Optional)

20

Estimated performance

Operation Type	Read	Write
Sustained random IOPS limit	6	30
Sustained throughput limit (MB/s)	2.4	1.8

Encryption

Automatic (recommended)


Create

Cancel

Equivalent [REST](#) or [command line](#)

Snapshotting an instance

And once you have a boot disk, you can create a VM from it:

 Compute Engine

VM instances

Instance groups

Instance templates

Disks

Snapshots

Images

Metadata

Health checks

Zones

← Create a new instance


Name ⓘ
myfirstvm-1

Zone ⓘ
us-central1-b

Machine type

1 vCPU3.75 GB memoryCustomize

Boot disk ⓘ

 20 GB standard persistent disk
Existing disk
firstvm-boot-disk

Change

Replicating VMs

A boot disk can only be attached in RW mode to a single instance. To attach a boot disk in read-only mode to many instances, use

```
$ gcloud compute instances create vm-{1..3} --zone us-central1-b \
  --disk name=firstvm-boot-disk,boot=yes,mode=ro
```

A read-only boot disk is a pain! This is how we get identical snapshots with read/write boot disks:

```
$ gcloud compute disks create vm-{1..5} \
  --source-snapshot "vm-snapshot" --zone us-central1-b
$ for i in {1..5}; do gcloud compute instances create vm-$i \
  --zone=us-central1-b --disk name=vm-$i,mode=rw,boot=yes; done
```

Exercise

You created an n1-standard-1 VM to serve a website. However, since this VM is handling the load with very little CPU usage, you realize you can save money by using a shared-CPU instance.

Exercise: Use persistent boot disks to migrate your server from an n1-standard-1 instance to an f1-micro instance.

Exercise

Step 1: Create the standard instance:

```
$ gcloud compute instances create myserver \
  --zone=us-central1-b --machine-type=n1-standard-1 \
  --tags=http-server
```

Created [<https://www.googleapis.com/compute/v1/projects/learncloud-118>]

NAME	ZONE	MACHINE_TYPE	PREEMPTIBLE	INTERNAL_IP	EXTERNAL_I
myserver	us-central1-b	n1-standard-1		10.240.0.3	130.211.12

Exercise

Step 2: Start an nginx-server on this instance:

```
$ gcloud compute ssh myserver --zone=us-central1-b
myserver$ sudo apt-get update && sudo apt-get install -y nginx
myserver$ sudo chmod a+rw /var/www/html/index.nginx-debian.html
myserver$ echo "Hola!" > /var/www/html/index.nginx-debian.html
myserver$ exit
$ curl 130.211.120.11 # Make sure it works!
Hola!
```

Exercise

Step 3: Delete your instance, but keep the boot disk!

```
$ gcloud compute instances delete myserver \
  --keep-disks boot --zone=us-central1-b
```

```
$ gcloud compute disks list
```

NAME	ZONE	SIZE_GB	TYPE	STATUS
myserver	us-central1-b	10	pd-standard	READY

Exercise

Step 4: Create a new instance from the boot disk:

```
$ gcloud compute instances create smaller-server \
--zone=us-central1-b --machine-type=f1-micro \
--tags=http-server --disk name=myserver,boot=yes,mode=rw
Created [https://www.googleapis.com/compute/v1/projects/learncloud-118
NAME          ZONE          MACHINE_TYPE  INTERNAL_IP  EXTERNAL_IP    S
smaller-server us-central1-b f1-micro      10.240.0.3   130.211.120.111 R
$ curl 130.211.120.111
Hola!
```

Editing Firewalls

Google cloud has a set of default firewall rules that you might like to edit:

Addresses				
10.240.0.0/16				
Gateway				
10.240.0.1				
Firewall rules				
Add firewall rule		Delete		
<input type="checkbox"/>	Name ^	Source tag / IP range	Allowed protocols / ports	Target tags
<input type="checkbox"/>	default-allow-http	0.0.0.0/0	tcp:80	http-server
<input type="checkbox"/>	default-allow-https	0.0.0.0/0	tcp:443	https-server
<input type="checkbox"/>	default-allow-icmp	0.0.0.0/0	icmp	Apply to all targets
<input type="checkbox"/>	default-allow-internal	10.240.0.0/16	tcp:0-65535; udp:0-65535; icmp	Apply to all targets
<input type="checkbox"/>	default-allow-rdp	0.0.0.0/0	tcp:3389	Apply to all targets
<input type="checkbox"/>	default-allow-ssh	0.0.0.0/0	tcp:22	Apply to all targets

Note that all ports are available via the internal ip addresses 10.240.0.0/16

```
$ gcloud compute firewall-rules list
```

NAME	NETWORK	SRC_RANGES	RULES
default-allow-http	default	0.0.0.0/0	tcp:80
default-allow-https	default	0.0.0.0/0	tcp:443
default-allow-icmp	default	0.0.0.0/0	icmp
default-allow-internal	default	10.240.0.0/16	tcp:0-65535,udp:0-65535,icmp
default-allow-ssh	default	0.0.0.0/0	tcp:22

Editing Firewalls

The default rules are of course reflected in the port scan:

```
$ nmap -p 0-10000 104.154.88.253
```

```
Starting Nmap 6.47 ( http://nmap.org ) at 2016-01-08 17:43 CST
Nmap scan report for 253.88.154.104.bc.googleusercontent.com (104.1
Host is up (0.040s latency).
Not shown: 9997 filtered ports
PORT      STATE SERVICE
22/tcp    open  ssh
80/tcp    closed http
443/tcp   closed https
3389/tcp   closed ms-wbt-server
```

Editing firewalls

This allows traffic from AMQP:

Firewall rules

Description (Optional)

Network ?

Source filter ?

Allowed protocols and ports ?

Target tags (Optional) ?

Editing firewalls

Command line to allow AMQP on port 5672 from any source IP, use the command

```
$ gcloud compute firewall-rules \
  create default-allow-amqp --allow tcp:5672 \
  --target-tags rabbit-server
```

Created [<https://www.googleapis.com/compute/v1/projects/learncloud-118>]

NAME	NETWORK	SRC_RANGES	RULES	SRC_TAGS	TARGET_TAGS
default-allow-amqp	default	0.0.0.0/0	tcp:5672	rabbit-server	

Creating a firewall rule applies it to all live instances in the project, unless you give it a target-tag.

Target Tags

Target tags give us an idea of what our instance is going to be used for. This tag allows traffic on port 80/443:

```
$ gcloud compute instances create webserver \
  --zone us-central1-b --tags http-server,https-server
```

Our custom amqp firewall rule will be applied to this server:

```
$ gcloud compute instanecs create rabbit \
  --zone us-central1-b --tags rabbit-server
```

Networking

Each project supports a $2^{16} = 65,536$ address internal network:

```
$ gcloud compute networks list
```

NAME	IPV4_RANGE	GATEWAY_IPV4
default	10.240.0.0/16	10.240.0.1

Load balancing

Let's set up a couple servers to learn how to load balance. First, our startup script.sh:

```
#!/bin/bash
sudo apt-get update
sudo apt-get install -y nginx emacs
sudo chmod a+rw /var/www/html/index.nginx-debian.html
sudo echo "Hello from server 1!" > /var/www/html/index.nginx-debian.htm
```

Load balancing

Now create your servers with this startup script:

```
$ gcloud compute instances create server-1 \
  --tags http-server --zone us-central1-a \
  --metadata-from-file startup-script=script.sh
```

Edit the file so that it tells us what server it comes from; do the same for server-2 and server-3, editing the script (“Hello from server 2!”, so on).

Load Balancing

Now let's set up a "target-pool", a group of server that responds to requests:

```
$ gcloud compute target-pools create \
    mypool --region us-central1
$ gcloud compute target-pools add-instances \
    mypool --instances server-{1..3} --zone us-central1-a
$ gcloud compute forwarding-rules create \
    frontend-forwarding --target-pool mypool --region us-central1
```

NAME	REGION	IP_ADDRESS	IP_PROTOCOL	TARGET
frontends	us-central1	104.154.36.151	TCP	us-central1/targetPool

```
$ for i in {1..100}; do curl 104.154.36.151; done
Welcome to server-1!
Welcome to server-1!
Welcome to server-3!
Welcome to server-2!
Welcome to server-1!
```

Load Balancing

What happens if one of our servers stops?

```
$ for i in {1..100}; do curl 104.154.44.41; done
```

```
Hello from server-1!
```

```
Hello from server-3!
```

```
curl: (7) Failed to connect to 104.154.44.41 port 80: Connection refused
```

The load balancer doesn't know anything about our servers, so it can't tell if it should send traffic there or not. Let's educate our load balancer:

Educating your load balancer

We need our load balancer to know that the server on port 80 should be available. Therefore, we create a health-check for it:

```
$ gcloud compute http-health-checks create frontend-check
```

```
Created [https://www.googleapis.com/compute/v1/projects/learncloud-118
```

NAME	HOST	PORT	REQUEST_PATH
------	------	------	--------------

frontend-check		80	/
----------------	--	----	---

```
$ gcloud compute target-pools add-health-checks mypool \
  --http-health-check frontend-check --region=us-central1
```

```
Updated [https://www.googleapis.com/compute/v1/projects/learncloud-118
```


Load Balancer with Health Checks

Now we shouldn't see many failed requests:

```
$ for i in {1..100}; do curl 104.154.44.41; done
```

Hello from server-1!

Hello from server-3!

Hello from server-1!

Hello from server-1!

Hello from server-3!

Hello from server-3!

Hello from server-3!

Hello from server-1!

Hello from server-3!

Hello from server-1!

Hello from server-3!

Hello from server-1!


Hello from server-1!


Hello from server-3!


Hello from server-3!


Educating your load balancer


For those who don't like the command line, of course we can create the health check in the console:


 Compute Engine


 VM instances


 Instance groups


 Instance templates


 Disks


 Snapshots


 Images


 Metadata

 Health checks

 Zones

 Operations

 Quotas

 Settings

← Create a health check

Autohealing instance groups and load balancing use health checks to detect when an instance is unresponsive. [Learn more](#)

Name ⓘ

Description (Optional)

Protocol

HTTPS

Port ⓘ

Request path ⓘ

⌵ More

Health criteria

Define how health is determined: how often to check, how long to wait for a response, and how many successful or failed attempts are decisive.

Check interval ⓘ	<input type="text" value="5"/> seconds	Timeout ⓘ	<input type="text" value="5"/> seconds
Healthy threshold ⓘ	<input type="text" value="2"/> consecutive successes	Unhealthy threshold ⓘ	<input type="text" value="2"/> consecutive failures

Create

Cancel

Instance Templates

Before autoscaling, we need to create an *instance template*, which describes how big we want our nodes, what OS, and perhaps a startup script. Remember, these nodes must be identical.

Compute Engine

VM instances

Instance groups

Instance templates

Disks

Snapshots

Images

Metadata

Health checks

Zones

Operations

Quotas

Create a new instance template

Use an instance template to describe a VM instance once and then create groups of identical instances. [Learn more](#)


Name

server-template

Machine type

micro (1 shared vCPU) 0.6 GB memory [Customize](#)

Boot disk

 New 10 GB standard persistent disk
Image
Debian GNU/Linux 8.2 (jessie) [Change](#)

Firewall

Add tags and firewall rules to allow specific network traffic from the Internet

☒ Allow HTTP traffic

☒ Allow HTTPS traffic

\$4.49 per month estimated

Effective hourly rate \$0.006 (730 hours per month)

Show costs for location

US

[Details](#)

You can also use a custom image instead of an instance template; this is preferred for complex deployments.

Instance Templates

To create the instance template from the command line using a startup script, use:

```
$ gcloud compute instance-templates create my-instance-template \
  --machine-type=f1-micro --scopes=compute-rw,storage-full \
  --tags=http-server \
  --metadata-from-file startup-script=server_startup.sh
```

Question 1: Why didn't we need to specify a zone here?

Question 2: Why do we need to tag this with "http-server"?

Instance Groups

Once we have an instance template, we can create an autoscaling instance group:

Compute Engine	← Create a new instance group
VM instances	Use an instance group when configuring a load-balancing backend service or to group VM instances. Learn more
Instance groups	Name ⓘ server-group
Instance templates	Description (Optional) <div></div>
Disks	Zone ⓘ us-central1-b
Snapshots	Specify port name mapping (Optional) <div>Use instance template Select existing instances</div>
Images	Creation method Use a template to create a group of identical instances that can scale automatically. If you do not use a template, you must add and manage each member yourself. Learn more
Metadata	Instance template ⓘ server-template
Health checks	Autoscaling ⓘ On
Zones	Autoscale based on ⓘ For best results read Configuring autoscaling instance groups CPU usage
Operations	Target CPU usage ⓘ Scaling dynamically creates or deletes VMs to meet the group target. Learn more 60 %
Quotas	
Settings	

Instance Groups

Command-line for creating instance groups:

```
$ gcloud compute instance-groups managed create my-instance-group \
  --zone us-central1-b --template my-instance-template \
  --base-instance-name my-instance-group --size 1
```

Autoscaling an Instance Group

Command line:

```
$ gcloud compute instance-groups managed set-autoscaling "my-instance-g  
  --zone "us-central1-b" --cool-down-period "60" \  
  --max-num-replicas "4" --min-num-replicas "1" \  
  --target-cpu-utilization "0.6"
```

This should now start autoscaling based on load!

Zonal Data Disks

In addition to creation of boot disks, you can create zonal data disks:

```
$ gcloud compute disks create "data-disk" \  
  --size=200 --zone=us-central1-b
```

Created [<https://www.googleapis.com/compute/v1/projects/learncloud-118>]

NAME	ZONE	SIZE_GB	TYPE	STATUS
data-disk	us-central1-b	200	pd-standard	READY

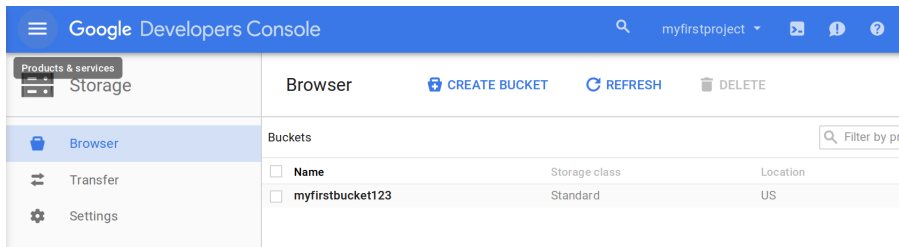
Zonal Data Disks

We can now mount this on an instance:

```
$ gcloud compute instances create datavm \
  --scopes=compute-rw --zone=us-central1-b
$ gcloud compute ssh datavm --zone=us-central1-b
datavm$ gcloud compute instances attach-disk datavm --disk data-disk --
datavm$ sudo mkdir -p /mnt/pd0
datavm$ ls -l /dev/disk/by-id/google*
datavm$ sudo /usr/share/google/safe_format_and_mount -m "mkfs.ext4 -F"
           /dev/disk/by-id/google-persistent-disk-1 /mnt/pd0
datavm$ df -k
...
```

Persistent Disks

If you need storage that is not chained to a particular zone, then you need to use “buckets”:



The screenshot shows the Google Developers Console interface. The top navigation bar is blue with the text "Google Developers Console" and a search icon. Below the navigation bar, the left sidebar shows "Products & services" with a "Storage" link. The main content area is titled "Browser" and contains a "CREATE BUCKET" button, a "REFRESH" button, and a "DELETE" button. Below these buttons is a table of buckets. The table has columns for "Name", "Storage class", and "Location". One bucket is listed: "myfirstbucket123" with a "Standard" storage class and "US" location. A search bar with the text "Filter by p" is visible on the right side of the table.

Name	Storage class	Location
myfirstbucket123	Standard	US

Persistent Disks

In order to manage persistent disks from the command line, use the gsutil utility:

```
$ pip install gsutil
$ gsutil mb gs://myfirstbucket
Creating gs://myfirstbucket/...
ServiceException: 409 Bucket myfirstbucket already exists.
$ gsutil mb gs://myfirstbucket123
Creating gs://myfirstbucket123/...
$ gsutil cp talk.log gs://myfirstbucket123
Copying file://talk.log [Content-Type=application/octet-stream]...
Uploading gs://myfirstbucket123/talk.log: 43
~/learn_gcloud$ gsutil ls
gs://myfirstbucket123/
$ gsutil ls -l gs://myfirstbucket123
44993 2016-01-09T21:19:30Z gs://myfirstbucket123/talk.log
TOTAL: 1 objects, 44993 bytes (43.94 KiB)
```

Bucket names need to be globally unique across all of gcloud!

Persistent Disks

- ▶ gcloud buckets are encrypted on disk
- ▶ gcloud buckets are accessible outside zones (i.e. have global scope)
- ▶ Uploads to gcloud buckets are atomic and considered successful once the info is stored in multiple datacenters.

Persistent Disks

Mounting a bucket to an instance takes a little work ([instructions](#)). First let's create a startup script called can-bucket.sh:

```
#!/bin/bash
sudo apt-get update
sudo apt-get install -y fuse emacs
sudo curl -L -O \
https://github.com/GoogleCloudPlatform/gcsfuse/releases/download/v0.15.
sudo dpkg --install gcsfuse_0.15.0_amd64.deb
```

Persistent Disks

Now let's create some instances which have permission to view our bucket:

```
$ gcloud compute instances create can-bucket-{1..3} \
  --zone=us-central1-b --scopes=storage-full \
  --metadata-from-file startup-script=can-bucket.sh
```

The “scopes=storage-full” gives our instances permission to mount our buckets.

Persistent Disks

```
$ gcloud compute ssh can-bucket-1 --zone=us-central1-b  
can-bucket-1$ mkdir mount_point  
can-bucket-1$ gcsfuse myfirstbucket123 mount_point  
can-bucket-1$ touch mount_point/file.txt
```

That'll do it! Try it on your other instances!

Persistent Disks

A few more gsutil examples:

```
$ gsutil rsync gs://myfirstbucket123/ 'pwd'  
Building synchronization state...  
Starting synchronization  
Copying gs://myfirstbucket123/talk.log...  
Downloading file:///home/NAThompson/talk.log:  
$ gsutil rsync 'pwd' gs://myfirstbucket123/foo  
Building synchronization state...  
Starting synchronization  
...  
$ gsutil du -h gs://myfirstbucket123  
43.94 KiB    gs://myfirstbucket123/talk.log
```

43.94

Service Accounts

What if I want my instance to be able to launch other instances? It can't login to my gmail account, so how can this be done? The answer is "service accounts". Service accounts can credential non-human agents and allow them to rent resources.

```
$ gcloud compute instances create withservice --scopes compute-rw --zone=us-central1-b
$ gcloud compute ssh withservice --zone=us-central1-b
wserviceacct$ gcloud compute instances create another-instance --zone=us-central1-b
Created [https://www.googleapis.com/compute/v1/projects/learncloud-118/instances/another-instance]
NAME                                ZONE                MACHINE_TYPE  PREEMPTIBLE  INTERNAL_IP  EXTERNAL_IP
another-instance  us-central1-b      n1-standard-1                10.240.0.5   10.240.0.5
```

However, without a service account scope, we get:

```
$ gcloud compute instances list
NAME  ZONE  MACHINE_TYPE  PREEMPTIBLE  INTERNAL_IP  EXTERNAL_IP  STATUS
ERROR: (gcloud.compute.instances.list) Some requests did not succeed:
- Insufficient Permission
```

Service Accounts

Note that even though I gave my instance access to a service account, it still can't read/write to my bucket:

```
withcomputeserviceacct$ gsutil ls gs://mybucket123  
AccessDeniedException: 403 Insufficient Permission
```

To fix this, we need to give our instance permission to read/write cloud storage:

```
$ gcloud compute instances create can-bucket \\  
    --zone=us-central1-b --scopes=compute-rw,storage-rw
```

Now gcsfuse will succeed in a read/write mount of our cloud storage.

Exercise

Why does the following command fail?

```
$ gcloud compute instances create somevm \  
--zone=us-central1-b \  
--metadata startup-script-url=gs://myfirstbucket123/gs_startup_script
```

Estimating Costs

gcloud gives a nice website for estimating costs:

<https://cloud.google.com/products/calculator>

References

- ▶ [Google Compute Engine](#), by Marc Cohen, Kathryn Hurley, and Paul Newson
- ▶ [Building Your Next Big Thing With Google Cloud Platform](#), by Jose Gonzales and S. Krishnan.